

000  
001  
002

054

003 

# High-Resolution Pose Transfer via Progressive Training and Pose Disentangling

055

004  
005  
006  
007  
008  
009  
010  
011

056

Anonymous CVPR submission

057

012 Paper ID \*\*\*\*  
013

058

## Abstract

059

This paper proposes a novel High-Resolution Pose Transfer Network (HPN) which transfers an arbitrary target pose with unprecedented image resolution ( $1024^2$ ) to a reference person, given only an image of the same person with the target pose. Our HPN framework utilizes dense local descriptors to refine local details, which are trained progressively in a coarse-to-fine manner to produce the high-resolution output to faithfully preserve the complex appearance of garment textures and geometry, while transferring seamlessly the target pose including those with self-occlusion. Our progressive encoder-decoder architecture can disentangle pose from appearance inherent the input image at multiple scales. Extensive experimental results on Human3.6M [14], DeepFashion [26] and our dataset collected from YouTube show that our model produces high-quality images, which can be further utilized in useful applications such as high-quality garment transfer between different persons and pose-guided person video generation.

060

061

062

063

064

065

066

067

068

069

070

071

072

073

074

075

076

077

078

079

080

081

082

083

084

085

086

087

088

089

090

091

092

093

094

095

096

097

098

099

100

101

102

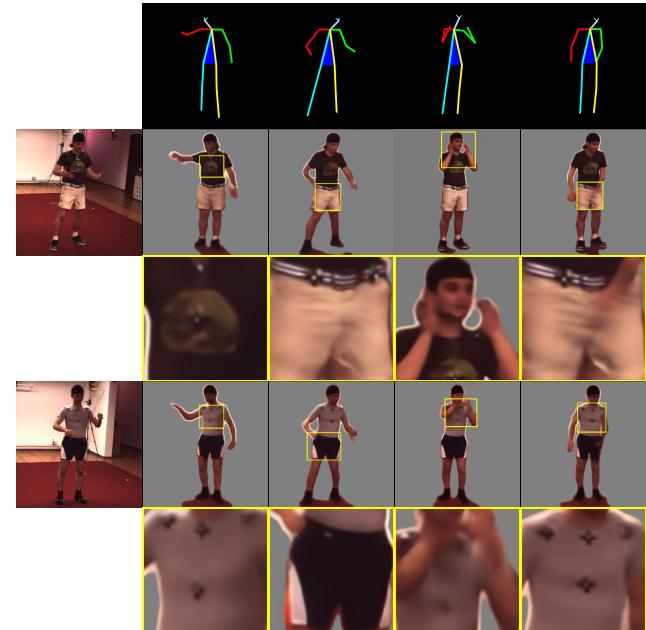
103

104

105

106

107



**Figure 1. Transfer results.** Given a reference image (leftmost column) and target poses (top row) as input which contains self-occlusion with complex appearance in texture and geometry, our HPN can transfer the target pose to the reference person in high resolution preserving high level of details. (a), (b) demonstrate the effects of dis-occlusion and (c), (d) demonstrate the effects of transferring to other self-occluded poses with zoom-in shown in detail.

034  
035

## 1. Introduction

036

Learning 3D information inherent in the 2D image domain is a fundamental problem in computer vision. This problem is fundamental in many computer vision tasks such as scene understanding [24], instance segmentation [5, 23] and action recognition [38, 8], to name a few, while remaining a major challenge for deep neural network learning. The challenge lies in the fact that images are 2D projections of the corresponding 3D world where objects can undergo complex deformation and occlusion.

037

This paper focuses on images of humans, whose different poses introduce complex non-rigid deformation and self-occlusion. Specifically, given a reference image of a person and another image of the same person in a target pose, our method seamlessly transfers the target pose to the reference person while preserving high-quality garment texture of the reference person, and at the same time hallucinating realistically their complex appearance under the target pose, see Figure 1. Note that the network must not

038

039

040

041

042

043

044

045

046

047

048

049

050

051

052

053

only moves the corresponding body parts to match the target pose, but also realistically inpaint or hallucinate exposed body/garment parts unseen in the input due to occlusion. To this end, the network must learn to disentangle the structure and appearance of the reference person from the given image. This is particularly challenging for human images due to the non-rigid nature of 3D human body, and complex texture and geometry distortion on the 3D garment worn by humans.

Many recent works seem to provide a plausible solution to our human pose transfer task. Conditional Generative

108 Adversarial Networks (GANs) [15], for example, have been  
109 exploited to effectively solve similar tasks such as generating  
110 MNIST [22] digits given labels. In particular, they  
111 generate sharp and realistic images based on certain pre-  
112 conditions by minimizing an adversarial loss. However,  
113 they can only generate images in accordance to the training  
114 distribution, but fail to reconstruct or hallucinate un-  
115 seen details, and thus are not applicable to our task which  
116 may involve dis-occlusion to reveal unseen details. Though  
117 the other plausible generative model – the Variational Au-  
118 toencoders (VAEs) [20] – can generate results complying  
119 a given reference image, they may not adequately preserve  
120 high-quality details in the reference image, due to the fact  
121 that the related method maximizes only a lower bound.  
122

123 To address these limitations, we propose the High-  
124 resolution Pose Transfer Network (HPN), which is effec-  
125 tive in disentangling structure and appearance information  
126 inherent in a given reference image, and faithfully transfers  
127 the original appearance of the person according to the tar-  
128 get pose representation. Figure 2 gives an overview of HPN.  
129 Specifically, we inject the target pose representation into the  
130 bottleneck of our encoder-decoder architecture for disen-  
131 tangling pose and appearance. Then we adopt local descriptors  
132 on image regions to encourage the network to learn more  
133 details for enhancing generation quality. Furthermore,  
134 progressive growth is employed on both encoder and decoder  
135 sides to increase output resolution. During training, per-  
136 ceptual loss [17], globally and locally applied around the  
137 regions of local descriptors, is used to compare between the  
138 output and the ground truth in the target pose.  
139

140 To validate our approach, we conduct extensive exper-  
141 iments on Human3.6M [14], DeepFashion [26] and our  
142 dataset collected from YouTube, and our results show that  
143 the HPN outperforms current state-of-the-art generative  
144 models. We also apply HPN to other applications such as  
145 high-quality garment transfer and pose-guided person video  
146 generation. Our results demonstrate the high potential of  
147 HPN in many challenging tasks.  
148

149 Our contribution is three-fold: 1) to disentangle struc-  
150 ture and appearance inherent in a given reference image, we  
151 propose a new encoder-decoder architecture that success-  
152 fully enables seamless human pose transfer; 2) we propose  
153 novel local descriptors to enhance the generation quality  
154 and local details; 3) we apply progressive training to our au-  
155 toencoder architecture to achieve outputs of unprecedented  
156 high resolution ( $1024^2$ ). To our knowledge, this is the first  
157 progressive, deep autoencoder transfer network that can re-  
158 alistically hallucinate in high resolution at the target pose  
159 the complex appearance of the worn garment, including the  
160 portion that was previously occluded in the reference image.  
161

## 2. Related Work

162  
163 **Conditional image generation** Generative models includ-  
164 ing Variational Autoencoders [20] (VAEs) and GANs[9]  
165 had demonstrated great power in image generation.  
166

167 Human pose transfer is closely related to the problem  
168 of conditional image synthesis as it requires a target pose  
169 as an output constraint. Zhao *et al.* [42] integrated GANs  
170 and other inference models to generate images of persons  
171 in various clothing styles from multiple views. Reed *et*  
172 *al.* [33] proposed a conditional generative model that used  
173 pose and text as conditions to generate images. Lassner *et*  
174 *al.* [21] also presented a generative model based on human  
175 pose that could generate realistic images conditioning on  
176 clothing segmentation.

177 Numerous researchers [36, 31, 40, 41, 45] introduced  
178 their respective methods to allow more control on the ap-  
179 pearances of the generated images in generative processes  
180 by providing different intermediate information such as la-  
181 bels and texts. Models such as ConditionalGAN [15] and  
182 CycleGAN [44] also demonstrated their efficacy in image-  
183 to-image translation. Yet, compared with our feedforward  
184 autoencoder, they and GANs in general are relatively more  
185 difficult to train, which often cannot faithfully transfer in-  
186 tricate patterns and textures from reference images.  
187

188 In general, it is difficult for the above methods to sim-  
189 taneously encode different factors such as pose and appear-  
190 ance. To transfer the pose-invariant human appearance, disen-  
191 tangling pose and appearance from reference images is an  
192 essential step. Many previous studies [3, 4] attempted to use  
193 GANs [9] and autoencoders [1] to disentangle such factors,  
194 including writing styles from character identities. Recently,  
195 Tran *et al.* [37] proposed DRGAN, which can disentangle  
196 pose from identity by learning the representation of human  
197 face followed by synthesizing the face with preserved iden-  
198 tity at the target pose.

199 **Pose transfer** There has been much work on pose transfer.  
200 Some approaches for pose transfer [6, 28] used encoder-  
201 decoders to attempt disentangling the pose and appearance  
202 of the input image to perform pose transfer. Esser *et al.* [7]  
203 explored a variational U-Net [30] on transferring the pose  
204 of a reference image invariant with its appearance. The  
205 PG<sup>2</sup> [27] was a more related work that aims at generating  
206 images of a subject in various poses based on an image of  
207 that person and one novel pose. Combining GANs and au-  
208 toencoders, PG<sup>2</sup> was trained through an encoder-decoder  
209 network followed by a refinement network given the pose  
210 and person image as input. Siarohin *et al.* [34] proposed a  
211 generative model similar to PG<sup>2</sup>, which added a discrimi-  
212 nator at the end of the autoencoder to help generate realistic  
213 images. Instead of using a discriminator, the pose trans-  
214 fer network presented by Natalia *et al.* [29] attempted to  
215 produce the seamless result by blending the synthesized im-

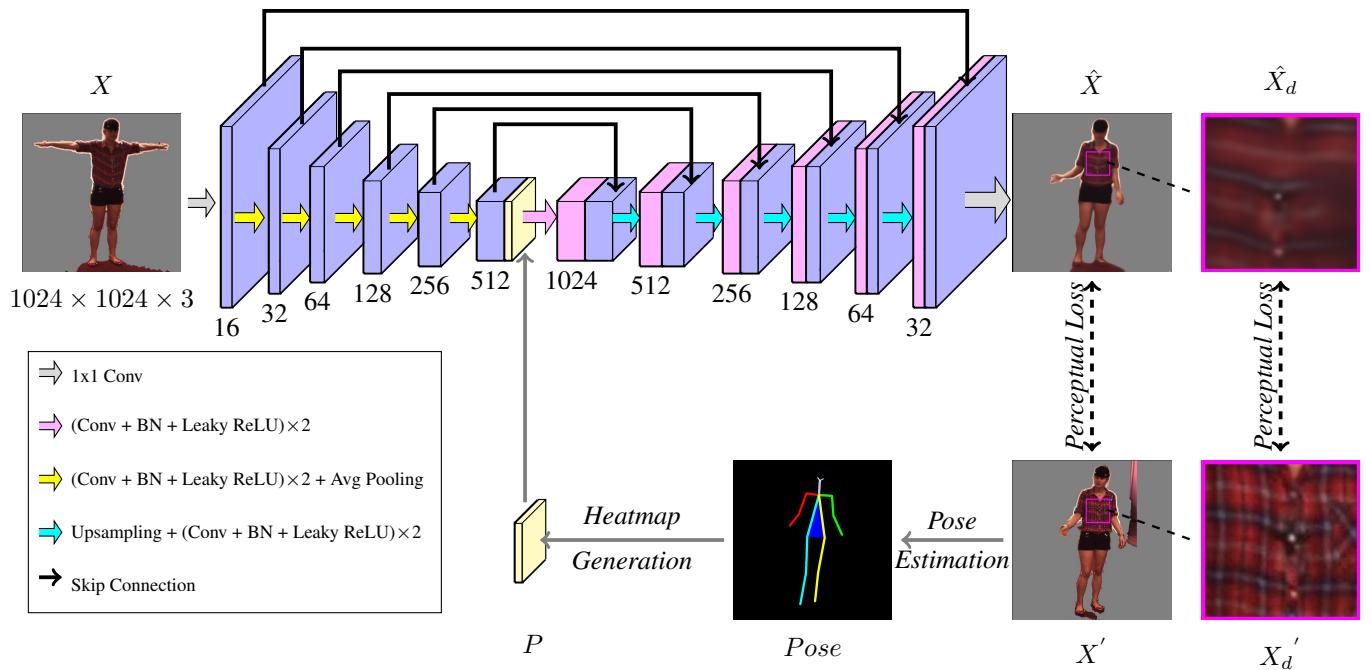


Figure 2. **Network Architectures.** The reference image  $X$  is first passed through an encoder to generate a latent representation. In the lower branch, 18 keypoints are estimated from the ground truth image  $X'$  to produce an explicit pose representation  $P$ .  $P$  is then concatenated with the latent representation, which is further decoded into the output  $\hat{X}$ . Global perceptual loss is enforced between  $X'$  and  $\hat{X}$ . To improve local details, local perceptual loss is also enforced on the corresponding local regions ( $X_d'$ ,  $\hat{X}_d$ ), indicated by the bounding boxes.

age and warped image through end-to-end training. Though not aiming at transferring human pose, the landmark learning network recently proposed by Jakab *et al.* [16] actually demonstrated acceptable results on pose transferring. This is achieved by using a simple encoder-decoder network with the learning landmarks concatenated in an intermediate representation.

In general, comparing to [27, 34], our method does not use sophisticated GANs which may introduce unstabilizing factors to the training process. More importantly, although [27, 34, 29, 16] performed well on changing pose at low-resolution ( $128 \times 128$ ) reference images while keeping their rough identity, they could not preserve but significantly blur complex textures after pose transfer.

**Progressive training** In the generative model, producing high-resolution and high-quality results is difficult since the training process becomes unstable and hard to converge as the output dimension increases. Recently, Tero *et al.* [18] proposed a progressive training methodology for generative adversarial networks to generate high-quality results. They started training from low resolution and added layers to the model progressively to obtain satisfactory high-resolution results. Ari *et al.* [13] also introduced an progressive architecture for autoencoder to encode and reconstruct high-

quality images (up to  $256^2$ ). They focused on how to train the autoencoder progressively for image reconstruction and image generation from the random sample while our goal is conditional image generation for even higher resolution output ( $1024^2$ ).

### 3. Method

Our goal is to transfer the pose of a reference person to a given target pose with high quality. This task is achieved by the disentanglement of appearance and pose of the person in the reference image through an autoencoder architecture. Our network architecture is shown in Figure 2.

Specifically, given a reference image  $X$  of a person and another image  $X'$  of the same person which is in the target pose, we first extract the explicit pose representation  $P$  from  $X'$  using a state-of-the-art pose estimator (Sec 3.1). We then inject  $P$  into the autoencoder’s bottleneck by concatenating it with the deepest feature map generated by the encoder. Finally, the concatenated feature block is passed through a decoder to generate an image with the person in the target pose, denoted as  $\hat{X}$ . Reconstruction loss is enforced globally between  $X'$  and  $\hat{X}$  to enforce pose transfer learning (Section 3.1). To improve generation quality, we adopt novel local descriptors to refine output details. Local

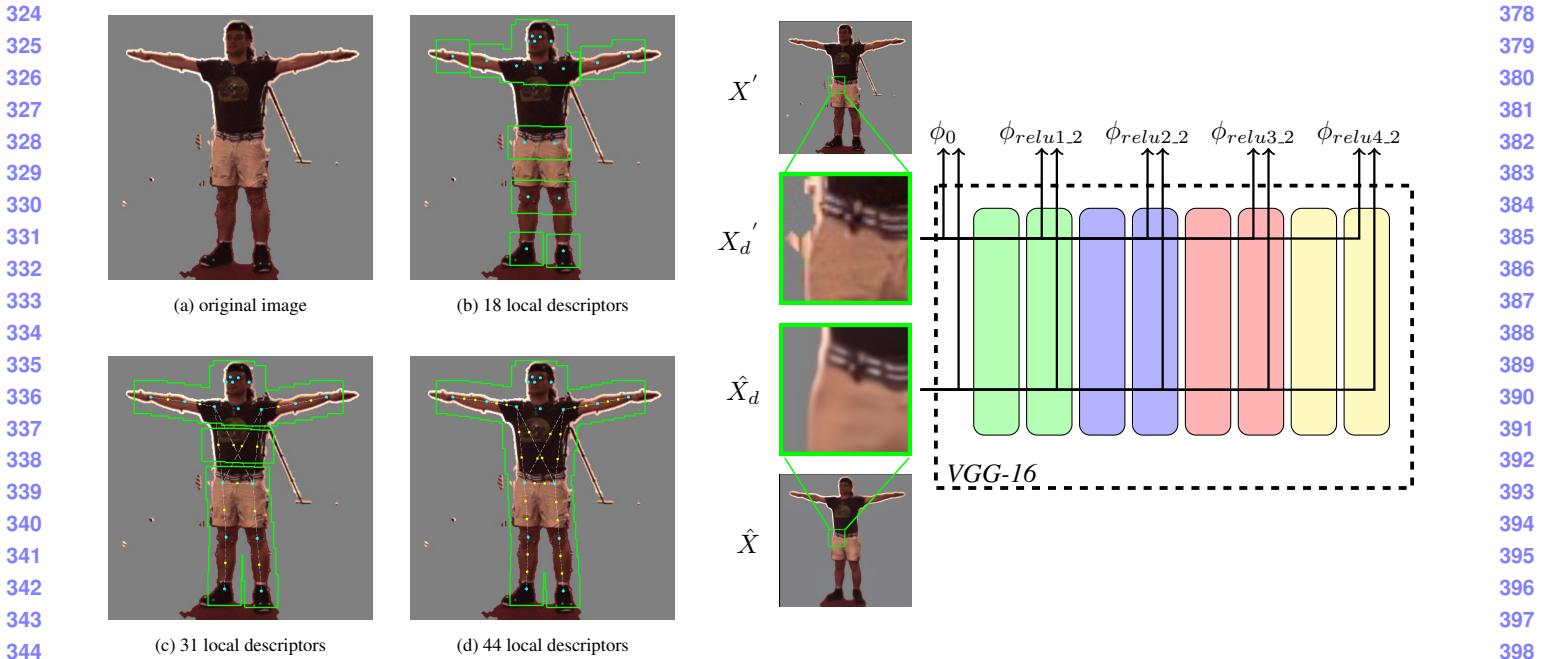


Figure 3. **Left:** the distribution and coverage of different numbers of local descriptors. Local descriptors are centered at the dots and their coverage is indicated by green bounding boxes. In particular, blue dots denote the 18 keypoints generated by a pose estimator and yellow dots denote the interpolated keypoints. Denser local descriptors introduce higher coverage of human body. **Right:** it demonstrates the mechanism of local loss back-propagation. Two corresponding local regions  $\hat{X}_d$  and  $X'_d$  are respectively cropped from generated image  $\hat{X}$  and ground truth image  $X'$ .  $\hat{X}_d$  and  $X'_d$  are then separately passed through a pre-trained VGG-16 to generate activations  $\phi$  at different layers  $l$ . A customized criterion  $C(\phi, \phi')$  measures the distances between corresponding activations  $\phi$ . Local descriptors intensify local loss back-propagation and thus enhance local details: see the sharper wrinkles and belt depicted in  $X'_d$ . Figure is best viewed online.

descriptors are applied under the guidance of keypoint locations from the pose estimator (Section 3.2). The same reconstruction loss is enforced locally at the corresponding regions described by local descriptors between  $X'$  and  $\hat{X}$ . To generate images in an unprecedented high resolution ( $1024^2$ ), the encoder and decoder are grown progressively as training proceeds (Section 3.3).

### 3.1. High-resolution Pose Transfer Network

**Pose representation** To represent human pose information in an explicit manner, we employ a state-of-the-art pose estimator [2], which gives the locations of 18 keypoints of a person in 2D coordinates. To let the network leverage the keypoint information effectively, these 18 keypoints are separately represented by a gaussian distribution map with a fixed standard deviation. Specifically, we denote each keypoint as  $k = 1, \dots, 18$  and their respective 2D coordinates as  $u(k)$ . Then the pose representation  $P$ , which is the concatenation of 18 gaussian distribution maps, is encoded as:

$$P(\mathbf{x}; k) = \exp\left(-\frac{1}{2\sigma^2} \|\mathbf{x} - u(k)\|^2\right) \quad (1)$$

The result is an explicit pose representation  $P \in \mathbb{R}^{H \times W \times 18}$  whose 18 maxima represent the locations of the 18 key-

points.  $P$  is then concatenated into the bottleneck of autoencoder.

**Autoencoder** The goal of the autoencoder is to reconstruct  $\hat{X}$  in the target pose based on the appearance of the person in the reference image  $X$  and the pose representation  $P$  extracted from the same person in the ground truth image  $X'$ , as shown in Figure 2. Since  $P$  contains no appearance information, the network is forced to utilize the appearance information in  $X$ . Furthermore, we add skip connections similar to that in a U-Net [30] to enable smoother gradient flow along the autoencoder. Then we adopt reconstruction loss between output  $\hat{X}$  and ground truth image  $X'$  to encourage the network to generate appropriate appearance which matches the pose of the person in  $X'$ .

**Perceptual loss** The design of reconstruction error is critical for good performance. Since it is hard for the network to learn a pixel-to-pixel mapping only from  $X$  due to the inherent pose and appearance variation, we encourage the network to also learn high-level semantic meanings during training, which is pivotal for pose and appearance disentanglement. Inspired by recent excellent practices [17], we adopt perceptual loss as the reconstruction loss between  $X'$  and  $\hat{X}$ . Apart from comparing only the raw pixel values, perceptual loss passes the output and the ground truth im-

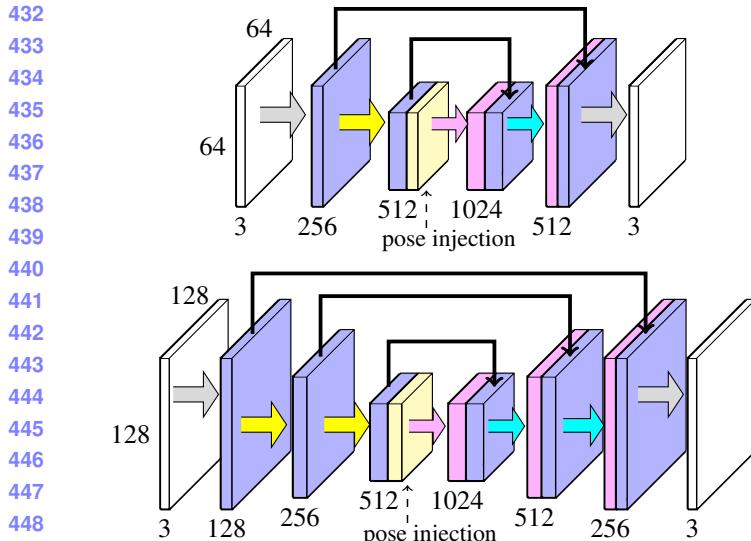


Figure 4. **Progressive training.** The bottleneck size of the autoencoder is  $32 \times 32$ . We start from a low spacial resolution of  $64 \times 64$  pixels and incrementally add layers to encoder and decoder as training proceeds until we reach the ultimate resolution of  $1024 \times 1024$ . All existing layers remain trainable throughout the process. Here we illustrate a snapshot when the network increases its resolution from  $64 \times 64$  to  $128 \times 128$ . During this transition, a new convolution block [ $(\text{Conv} + \text{BN} + \text{Leaky ReLU}) \times 2$ ] with corresponding up-sampling or down-sampling layer is introduced to encoder and decoder respectively.  $1 \times 1$  convolution layer used to project RGB channels to/from feature space is also replaced by a new one that fits the network.

ages individually through a pre-trained deep network and compares the activations extracted from multiple layers inside the network. This process enables the network to better learn the disentanglement of appearance and pose and alleviates overfitting. Specifically, we define perceptual loss as:

$$L(X', \hat{X}) = \sum_l C(\phi_l(X'), \phi_l(\hat{X})) \quad (2)$$

where  $\phi(x)$  is a pre-trained network, such as VGG-16 [35], and  $\phi_l$  denotes the activation of the  $l^{th}$  layer of  $\phi(x)$ . Different from common practices which use  $L_2$  loss as the criterion to evaluate  $\hat{X}$ , we customize the criterion  $C(\phi, \phi')$  to accelerate network convergence. Since  $L_2$  loss has an optimal solution while  $L_1$  loss enforces sharper output but is less stable, we designate  $C(\phi, \phi')$  as  $L_2$  loss in the first half of the training process within each resolution level and  $L_1$  loss in the second half. This practice enables stable convergence as well as high generation quality.

### 3.2. Local descriptors

The adoption of perceptual loss does not enforce sufficient preservation of local details. It is observed that sharp

garment textures cannot be well preserved under the restriction of global perceptual loss only, as shown in Figure 3. To address this limitation, we introduce novel local descriptors which enable generation of high-quality images. Local descriptors describe a set of regions telling the network where to focus and concentrate loss back-propagation. The locations of local descriptors are guided by the pose keypoints produced by the pose estimator. To ensure appropriate detail refinement and alleviate overfitting, the size of local regions is designed to be one-eighth of the input image resolution. The same reconstruction loss is applied locally between the corresponding regions in  $X'$  and  $\hat{X}$ .

Figure 3 shows the distribution and coverage of local descriptors. Since higher resolution generally requires more local details, we increase the number of local descriptors adopted by interpolating between existing keypoints as input image resolution grows. Denser overlapping local descriptors introduce more complete coverage of the body and thus help preserve details more faithfully.

Specifically, based on the 18 keypoints in  $X'$  produced by the pose estimator, a list of  $N$  local descriptors is generated, denoted as  $d = 1, \dots, N$ . Then two sets of fractional-sized regions centered at the location of each of  $N$  local descriptors are cropped from  $X'$  and  $\hat{X}$  respectively. Perceptual loss is enforced between corresponding local regions. The local loss  $L_{local}$  is formulated as the following:

$$L_{local}(X', \hat{X}) = \sum_{d=1}^N \sum_l C(\phi_l(X'_d), \phi_l(\hat{X}_d)) \quad (3)$$

where  $X'_d$  and  $\hat{X}_d$  denote the  $d^{th}$  region cropped from  $X'$  and  $\hat{X}$  respectively. Self-comparison between the model with and without local descriptors are shown in Figure 3. The significant improvement in image quality demonstrates promising enhancement introduced by local descriptors.

### 3.3. Progressive training of autoencoder

Apart from achieving high-quality image generation, we also aim at producing unprecedently high-resolution results ( $1024^2$ ). However, training the autoencoder in high resolution from scratch does not yield satisfactory results. Inspired by [13] which produces high-resolution results on CelebA-HQ dataset by introducing progressive training to GAN, we adopt a variation of progressive training which fits our setting of autoencoder with skip connections, as shown in Figure 4. Most importantly, instead of fading in a new convolution block to increase resolution using alpha blending, we train the new convolution block with skip connection from scratch, utilizing deeper convolution blocks trained in the previous stage as mature feature extractors. From our observation, this enables faster convergence of newly introduced blocks as well as utilization of skip connections to enhance generation quality. Self-comparison in

540 Figure 5 demonstrates substantial improvement brought by  
 541 progressive training on autoencoder.  
 542

### 543 3.4. Implementation details

544 We use the Adam [19] optimizer with a weight decay  
 545 of  $5 \times 10^{-4}$ . The initial learning rate is set to  $2 \times 10^{-4}$ .  
 546 We use  $\sigma = 3.2$  to generate the gaussian distribution for  
 547 pose representation. The autoencoder is trained progres-  
 548 sively starting from the resolution of  $64^2$  with bottleneck  
 549 shape of  $1024 \times 32^2$  and ending at the resolution of  $1024^2$ .  
 550 Within each convolution block, we use two contiguous sets  
 551 of  $3 \times 3$  convolution layer followed by batch normalization  
 552 [11] and leaky Relu with leakiness of 0.2. The number of  
 553 channels of feature maps is halved as spacial size doubles.  
 554 We downscale and upscale the feature maps using average  
 555 pooling and nearest neighbor interpolation respectively. We  
 556 use  $1 \times 1$  convolution to project the outermost feature maps  
 557 into RGB space and vice versa as in RGB back to feature  
 558 map. We use He’s initializer [12] to initialize the autoen-  
 559 coder. A total of 18 local descriptors are used for the  
 560 resolution of  $64^2$  and  $128^2$ . For  $256^2$  and  $512^2$ , we use 31  
 561 local descriptors by interpolating between keypoints pairs  
 562 and 44 local descriptors for  $1024^2$  through additional inter-  
 563 polations. For each resolution level, we train the network  
 564 for 700 thousand iterations.  
 565

566 Our final loss  $L$ , which is composed of both global loss  
 567  $L_{global}$  and local loss  $L_{local}$ , is formulated as the following:

$$\begin{aligned} L(X', \hat{X}) &= L_{global}(X', \hat{X}) + L_{local}(X', \hat{X}) \\ &= \sum_l C(\phi_l(X'), \phi_l(\hat{X})) \\ &\quad + \sum_{d=1}^N \sum_l C(\phi_l(X'_d), \phi_l(\hat{X}_d)) \end{aligned} \quad (4)$$

## 576 4. Experiments

577 To prove the advantages of the above proposed  
 578 method, we first conduct qualitative and quantitative self-  
 579 comparisons to validate the effectiveness of different com-  
 580 ponents of HPN, respectively local descriptors and progres-  
 581 sive training on autoencoders. We then demonstrate our  
 582 generalizability by showing the results produced by HPN  
 583 on various datasets, including human 3.6M [14], DeepFashion  
 584 [26] and a self-collected dataset from YouTube. We also  
 585 compare our performance on DeepFashion dataset with  
 586 previous work. Lastly, we show our potential to be further uti-  
 587 lized in real-world applications like high-quality garment  
 588 transfer between different persons and pose-guided person  
 589 video generation.

590 **Datasets** We train and test our model mainly on the Hu-  
 591 man3.6M dataset [14], which has 11 actors in total with dif-  
 592 ferent poses. The dataset provides ground truth 2D human  
 593

594 poses, backgrounds and human body bounding boxes. We  
 595 first subsample the sample videos at 3 frames per second  
 596 and obtain image frames with large pose variations. For  
 597 each image frame, we then subtract the background and re-  
 598 tain only the human foreground to reduce training noises.  
 599 We select ‘Posing’, ‘Greeting’ action classes for training,  
 600 and ‘Directions’ class for testing.

601 To test the generality of our method, we further train  
 602 and test on our self-collected youtube video datasets. The  
 603 datasets we collected has 10 dancing videos in total. All  
 604 of them have large pose variations. We subtract the back-  
 605 ground of this dataset using JPPNet [25] and subsample the  
 606 videos at 3 frames per second for training set as well as test-  
 607 ing set.

### 608 4.1. Self-comparison

**Local descriptors** Qualitative comparison in Figure 5  
 611 demonstrates the effectiveness of local descriptors. From  
 612 column (e), (g) and their corresponding zoom-in views, lo-  
 613 cal descriptors introduce improvement both in global co-  
 614 herency and local details compared to the baseline. And from  
 615 column (i), (k) and their corresponding zoom-in views, local  
 616 descriptors are still able to bring significant enhance-  
 617 ment to generation quality under progressive training.  
 618 In particular, the two stars in image (3, d) are faithfully pre-  
 619 served in result (3, l) but lost in result (3, j).

620 Figure ?? provides the validation of local descriptors  
 621 from another perspective, where the number of local de-  
 622 scriptors adopted by each model steadily increases. Specif-  
 623 ically, the four models are trained using 0, 18, 31, 44 local  
 624 descriptors respectively. Subtle but evident improvement  
 625 can be identified in the process of increasing the number of  
 626 overlapping local descriptors.

627 **Progressive training** The advantages of progressive train-  
 628 ing is also demonstrated through comparisons in Figure 5.  
 629 Column (g) and (k) with their corresponding zoom-in views  
 630 show the improvement for the models with local descrip-  
 631 tors, while Column (e) and (i) with zoom-in views show the  
 632 improvement for the models without local descriptors. In  
 633 particular, the garment texture in image (1, d) is faithfully  
 634 preserved in result (1, l) but lost in result (1, h). Even though  
 635 local descriptors help enhance local details, vanishing  
 636 gradient problem still persists in deep networks necessary for  
 637 high-resolution image generation. Progressive training en-  
 638 ables separate and progressive convergence of different lay-  
 639 ers in a deep network.

640 **Quantitative comparison** Image generation quality can be  
 641 hard to assess due to various standards. Here we adopt  
 642 Structural Similarity (SSIM) [39] as our main evaluation  
 643 metric. Due to the limitations of SSIM such as insensi-  
 644 tivity and distortion under-estimation near hard edge [32],  
 645 we also adopt a variation of SSIM, local-SSIM, to more  
 646 effectively evaluate local details. Instead of global evalua-

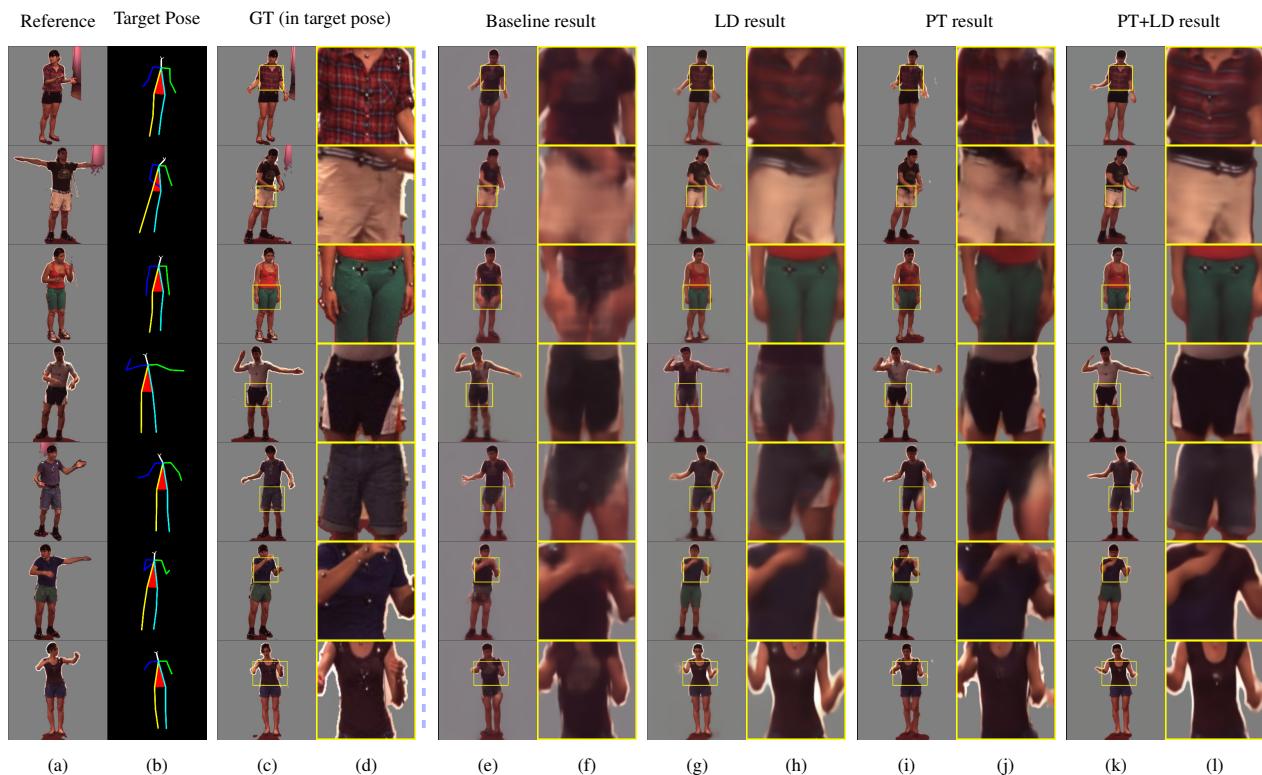


Figure 5. **Self-comparison results.** Test results on human 3.6M generated by Baseline (no local descriptors or progressive training), LD (with local descriptors), PT (with progressive training), LD + PT (with both local descriptors and progressive training) and their corresponding zoom-in views are provided. Local descriptors and progressive training each introduces considerable improvement in generation quality and produce the best result when combined. Our model also demonstrates robustness to the segmentation error introduced by Human 3.6M dataset.

ation performed by SSIM, local-SSIM operates on 44 corresponding local regions between the generated image and reference image. The 44 local regions correspond to the areas described by 44 local descriptors, where the highest coverage of human body is achieved.

Quantitative comparison between models under different settings are shown in Table 1. Either local descriptors or progressive training brings considerable enhancement in generation quality, with combination of the other further boost the result. Local SSIM more evidently reflects the improvement in the quality of local regions.

Table 1. Quantative self-comparison between different modes of our model.

Human3.6M		
Model	SSIM	local-SSIM
Baseline	0.909	0.699
LD	0.944	0.744
PT	0.953	0.759
LD+PT	<b>0.954</b>	<b>0.772</b>
Real Data	1.00	1.00

## 4.2. Youtube dataset results

## 4.3. Comparison with previous work

We compare our results with the current state-of-the-art method (**DSC**) on DeepFashion dataset. From the Figure 7

Table 2. Quantative comparison with previous work.

DeepFashion	
Model	SSIM
DSC	0.776
Ours	<b>0.806</b>
Real Data	1.00

## 4.4. Futher application

**Virtual try-on** Virtual try-on has seen great application potentials due to its convenience and reduction in cost. This problem involves the transfer of any garment with detailed and complex texture. While a recent approach [10] successfully preserve garment details and shapes, there still exist artifacts due to self-occlusions. With our HPN, we can tackle this problem with two steps. First, we transfer the

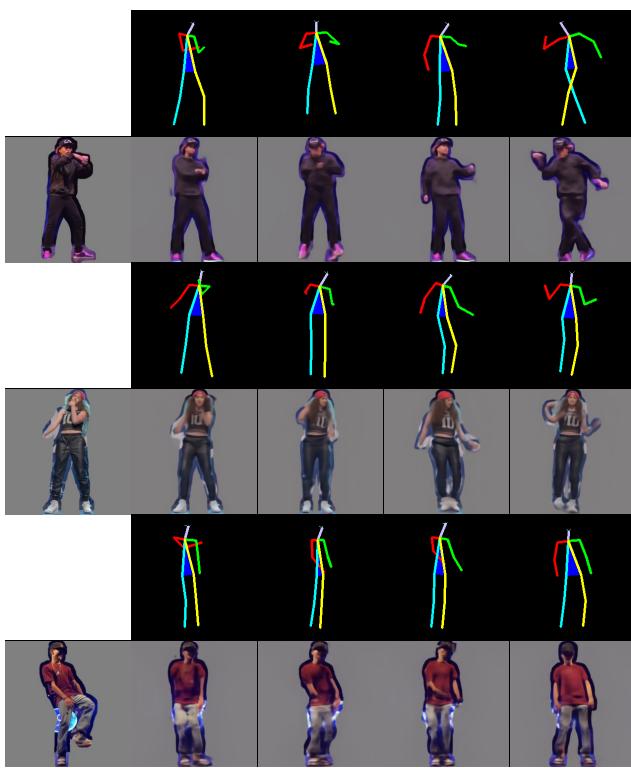


Figure 6. Youtube results.

image of the target person (with self-occlusion) into a pre-defined frontal pose (without occlusion). Then, we apply our appearance flow network to transfer the garment to that person.

Zhou *et al.* [43] first proposed the idea of appearance flows for view synthesis. Appearance flows are 2-D coordinate vectors describing how pixels in the input image could be used to reconstruct the image from the target viewpoint. They found that the images of the same target from different perspectives have high correlations with each other. Garment transfer is similar to view synthesis since the same garment piece in different human pose is also highly correlated. Inspired by their work, we propose an encoder-decoder framework to predict appearance flows for garment transfer. We first apply the JPPNet [25] to extract the mask of the garment in reference and target poses. The autoencoder will take the garment image, reference pose mask and target pose mask as input. Then the decoder outputs the appearance flows and yield the garment synthesized image through a bilinear sampling layer. We adopt perceptual loss between ground-truth image and synthesized image and update the layer’s weight through backpropagation.

**Video generation** Since our network can transfer the target pose to the reference person in high quality, it can be applied to human video generation given a sequence of target



Figure 7. Comparison with previous work.

poses.

## 5. Conclusion

## References

- [1] P. Baldi. Autoencoders, unsupervised learning and deep architectures. In *Proceedings of the 2011 International Conference on Unsupervised and Transfer Learning Workshop - Volume 27*, UTLW’11, pages 37–50. JMLR.org, 2011.
- [2] Z. Cao, T. Simon, S.-E. Wei, and Y. Sheikh. Realtime multi-person 2d pose estimation using part affinity fields. In *CVPR*, 2017.
- [3] X. Chen, Y. Duan, R. Houthooft, J. Schulman, I. Sutskever, and P. Abbeel. Infogan: Interpretable representation learning by information maximizing generative adversarial nets. In *NIPS*, 2016.

- |     |   |     |
|-----|---|-----|
| 864 | [4] B. Cheung, J. Livezey, A. Bansal, and B. Olshausen. Discovering hidden factors of variation in deep networks. In <i>ICLR workshop</i> , 2015.   | 918 |
| 865 |   | 919 |
| 866 |   | 920 |
| 867 | [5] J. Dai, K. He, and J. Sun. Instance-aware semantic segmentation via multi-task network cascades. In <i>CVPR</i> , 2016.   | 921 |
| 868 |   | 922 |
| 869 | [6] R. de Bem, A. Ghosh, T. Ajanthan, O. Miksik, N. Siddharth, and P. H. Torr. DGpose: Disentangled semi-supervised deep generative models for human body analysis. In <i>arXiv preprint arXiv:1804.06364</i> , 2018.                                       | 923 |
| 870 |   | 924 |
| 871 | [7] P. Esser, E. Sutter, and B. Ommer. A variational u-net for conditional appearance and shape generation. In <i>CVPR</i> , 2018.  | 925 |
| 872 |   | 926 |
| 873 | [8] R. Girdhar and D. Ramanan. Attentional pooling for action recognition. In <i>NIPS</i> , 2015.   | 927 |
| 874 |   | 928 |
| 875 | [9] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio. Generative adversarial nets. In <i>NIPS</i> , 2014.   | 929 |
| 876 |   | 930 |
| 877 | [10] X. Han, Z. Wu, Z. Wu, R. Yu, and L. S. Davis. Viton: An image-based virtual try-on network. In <i>CVPR</i> , 2018.   | 931 |
| 878 |   | 932 |
| 879 | [11] K. He, X. Zhang, S. Ren, and J. Sun. Batch normalization: Accelerating deep network training by reducing internal covariate shift. 2015.   | 933 |
| 880 |   | 934 |
| 881 | [12] K. He, X. Zhang, S. Ren, and J. Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. 2015.  | 935 |
| 882 |   | 936 |
| 883 | [13] A. Heljakka, A. Solin, , and J. Kannala. Pioneer networks: Progressively growing generative autoencoder. In <i>ACCV</i> , 2018.  | 937 |
| 884 |   | 938 |
| 885 | [14] C. Ionescu, D. Papava, V. Olaru, and C. Sminchisescu. Human3.6m: Large scale datasets and predictive methods for 3d human sensing in natural environments. <i>IEEE Transactions on Pattern Analysis and Machine Intelligence</i> , 36:1325–1339, 2014. | 939 |
| 886 |   | 940 |
| 887 | [15] P. Isola, J. Zhu, T. Zhou, and A. A. Efros. Image-to-image translation with conditional adversarial networks. In <i>CVPR</i> , 2017.   | 941 |
| 888 |   | 942 |
| 889 | [16] T. Jakab, A. Gupta, H. Bilen, and A. Vedaldi. Conditional image generation for learning the structure of visual objects. In <i>arXiv preprint arXiv:1806.07823</i> , 2018.   | 943 |
| 890 |   | 944 |
| 891 | [17] J. Johnson, A. Alahi, and L. Fei-Fei. Perceptual losses for real-time style transfer and super-resolution. In <i>ECCV</i> , 2016.  | 945 |
| 892 |   | 946 |
| 893 | [18] T. Karras, T. Aila, S. Laine, and J. Lehtinen. Progressive growing of gans for improved quality, stability, and variation. In <i>ICLR</i> , 2018.  | 947 |
| 894 |   | 948 |
| 895 | [19] D. P. Kingma and J. Ba. Adam: A method for stochastic optimization. <i>ICLR</i> , 2015.  | 949 |
| 896 |   | 950 |
| 897 | [20] D. P. Kingma and M. Welling. Auto-encoding variational bayes. In <i>ICLR</i> , 2014.   | 951 |
| 898 |   | 952 |
| 899 | [21] C. Lassner, G. Pons-Moll, and P. V. Gehler. A generative model of people in clothing. In <i>ICCV</i> , 2017.   | 953 |
| 900 |   | 954 |
| 901 | [22] Y. LeCun and C. Cortes. MNIST handwritten digit database. 2010.  | 955 |
| 902 |   | 956 |
| 903 | [23] G. Li, Y. Xie, L. Lin, and Y. Yu. Instance-level salient object segmentation. In <i>CVPR</i> , 2017.   | 957 |
| 904 |   | 958 |
| 905 | [24] L.-J. Li, R. Socher, and L. Fei-Fei. Towards total scene understanding: Classification, annotation and segmentation in an automatic framework. In <i>CVPR</i> , 2009.  | 959 |
| 906 |   | 960 |
| 907 |   | 961 |
| 908 |   | 962 |
| 909 |   | 963 |
| 910 |   | 964 |
| 911 |   | 965 |
| 912 |   | 966 |
| 913 |   | 967 |
| 914 |   | 968 |
| 915 |   | 969 |
| 916 |   | 970 |
| 917 |   | 971 |

972	[45] R. S. Zhu, S. Fidler, R. Urtasun, D. Lin, and C. Chen. Be	1026
973	your own prada: Fashion synthesis with structural coherence	1027
974	analysis. In <i>ICCV</i> , 2017.	1028
975		1029
976		1030
977		1031
978		1032
979		1033
980		1034
981		1035
982		1036
983		1037
984		1038
985		1039
986		1040
987		1041
988		1042
989		1043
990		1044
991		1045
992		1046
993		1047
994		1048
995		1049
996		1050
997		1051
998		1052
999		1053
1000		1054
1001		1055
1002		1056
1003		1057
1004		1058
1005		1059
1006		1060
1007		1061
1008		1062
1009		1063
1010		1064
1011		1065
1012		1066
1013		1067
1014		1068
1015		1069
1016		1070
1017		1071
1018		1072
1019		1073
1020		1074
1021		1075
1022		1076
1023		1077
1024		1078
1025		1079